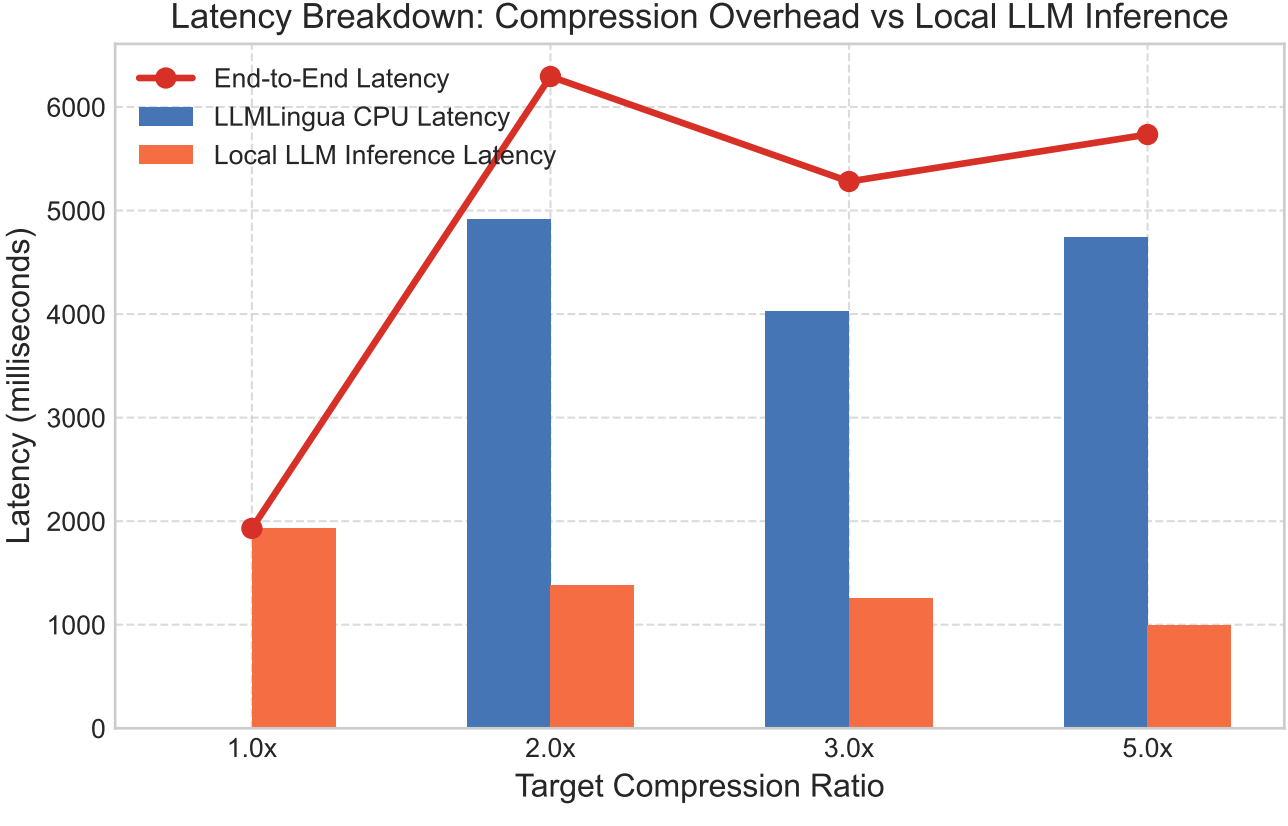
Latency Breakdown: Compression Overhead vs Local LLM Inference
End-to-End Latency
LLMLingua CPU Latency
Local LLM Inference Latency
Latency (milliseconds)
0
1000
2000
3000
4000
5000
6000
1.0x
2.0x
3.0x
5.0x
Target Compression Ratio